

1. Download all the data in this folder <https://drive.google.com/open?id=1Z4TyI7FcFVEx8qd14j09qxvxaqLSqoEu>. it contains two file both images and labels. The label file list the images and their categories in the following format:

path/to/the/image.tif,category

where the categories are numbered 0 to 15, in the following order:

- 0 letter**
- 1 form**
- 2 email**
- 3 handwritten**
- 4 advertisement**
- 5 scientific report**
- 6 scientific publication**
- 7 specification**
- 8 file folder**
- 9 news article**
- 10 budget**
- 11 invoice**
- 12 presentation**
- 13 questionnaire**
- 14 resume**
- 15 memo**

2. On this image data, you have to train 3 types of models as given below. You have to split the data into Train and Validation data.

3. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data.

or you can use this method also

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1> (<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>).

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c> (<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>).

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but y

ou have to use the same architecture what we are asking below.

5. Use tensorboard for every model and analyse your gradients. (you need to upload the screenshots for each model for evaluation)

Note: fit_generator() method will have problems with the tensorboard histograms, try to debug it, if you could not do use histograms=0 i.e don't include histograms, check the documentation of tensorboard for more information.

6. You can check about Transfer Learning in this link - <https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html> (<https://blog.keras.io/building-powerful-image-classification-models-using-very-little-data.html>).

```
In [2]: %tensorflow_version 2.x
```

```
In [3]: import numpy as np
import pandas as pd
import os
import tensorflow as tf
from sklearn.model_selection import train_test_split
```

```
In [4]: tf.__version__
```

```
Out[4]: '2.4.1'
```

```
In [10]: # goto this link: https://drive.google.com/uc?id=1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu
# click download anyway
# cancel the downloading...
# copy the link from CurlWget extension(if you don't have, install)
# then do => ! copied link from CurlWget extension
```

```
! wget --header="Host: doc-04-0k-docs.googleusercontent.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64; rv:68.0) Gecko/20100101 Firefox/68.0" https://doc-04-0k-docs.googleusercontent.com/docs/securesc/8iph74e76aalmppo83tp1i1i1vloj61q/55f323hme70a657df1j9sohnf4o8ub0h/1616462700000/00484516897554883881/08709549031353296611/1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu?e=download&authuser=0&nonce=2ld3kjsa27i4u&user=08709549031353296611&hash=lrp77usmuks0kc1foptibqeavm5m2cod (https://doc-04-0k-docs.googleusercontent.com/docs/securesc/8iph74e76aalmppo83tp1i1i1vloj61q/55f323hme70a657df1j9sohnf4o8ub0h/1616462700000/00484516897554883881/08709549031353296611/1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu?e=download&authuser=0&nonce=2ld3kjsa27i4u&user=08709549031353296611&hash=lrp77usmuks0kc1foptibqeavm5m2cod)
```

Resolving doc-04-0k-docs.googleusercontent.com (doc-04-0k-docs.googleusercontent.com)... 172.217.15.97, 2607:f8b0:4004:811::2001

Connecting to doc-04-0k-docs.googleusercontent.com (doc-04-0k-docs.googleusercontent.com)|172.217.15.97|:443... connected.

HTTP request sent, awaiting response... 200 OK

Length: unspecified [application/rar]

Saving to: 'rvl-cdip.rar'

rvl-cdip.rar [<=>] 4.34G 44.5MB/s in 87s

2021-03-23 01:27:05 (51.4 MB/s) - 'rvl-cdip.rar' saved [4660541790]

```
In [11]: !pip install patool
import patoolib
patoolib.extract_archive('rvl-cdip.rar')
```

Requirement already satisfied: patool in /usr/local/lib/python3.7/dist-packages (1.12)

patool: Extracting rvl-cdip.rar ...

patool: running /usr/bin/unrar x -- /content/rvl-cdip.rar

patool: with cwd='./Unpack_i_c7lxoo'

patool: ... rvl-cdip.rar extracted to `rvl-cdip' (multiple files in root).

```
Out[11]: 'rvl-cdip'
```

```
In [12]: # !gdown --id 1Z4TyI7FcFVEx8qdl4j09qxvxaqLSqoEu
# get_ipython().system_raw("unrar x rvl-cdip.rar")
```

```
In [13]: os.listdir('/content')
```

```
Out[13]: ['.config', 'rvl-cdip.rar', 'rvl-cdip', 'sample_data']
```

```
In [14]: dir_path = '/content/rvl-cdip/data_final'
```

```
In [15]: # reading label dataframe  
imagelabel = pd.read_csv('/content/rvl-cdip/labels_final.csv', dtype=str)  
imagelabel.head()
```

```
Out[15]:
```

	path	label
0	imagesv/v/o/h/voh71d00/509132755+-2755.tif	3
1	imagesl/l/x/t/lxt19d00/502213303.tif	3
2	imagesx/x/e/d/xed05a00/2075325674.tif	2
3	imageso/o/j/b/ojb60d00/517511301+-1301.tif	3
4	imagesq/q/z/k/qzk17e00/2031320195.tif	7

Let's see the distribution of the dataset

```
In [16]: class_dist = dict(imagelabel['label'].value_counts())
```

```
In [17]: class_dist
```


```
Out[17]: {'0': 3016,  
          '1': 2994,  
          '10': 3002,  
          '11': 2992,  
          '12': 3006,  
          '13': 3007,  
          '14': 3006,  
          '15': 2996,  
          '2': 2993,  
          '3': 3005,  
          '4': 2994,  
          '5': 2999,  
          '6': 2985,  
          '7': 3000,  
          '8': 3003,  
          '9': 3002}
```

```
In [18]: import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [19]: sns.barplot( list(class_dist.keys()), list(class_dist.values()) )
plt.title('class distribution of the dataset')
plt.xlabel('classes/labels')
plt.ylabel('No. of occurences')
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variables as keyword args: x, y. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.

FutureWarning



Observation:

1. Data is balanced. So, there is no need to balance the class.

```
In [19]:
```

```
In [20]: # imagelabel['label'] = pd.Series(map(lambda x: str(x), imagelabel.label))
```

```
In [21]: imagelabel.label.iloc[0]
```

```
Out[21]: '3'
```

```
In [22]: # an example of data point
import matplotlib.pyplot as plt
plt.figure(figsize=(10,10))
img_index = 2340
img = plt.imread(os.path.join(dir_path, imagelabel.path.iloc[img_index]))
plt.imshow(img)
plt.title(imagelabel.label.iloc[img_index])
```

Out[22]: Text(0.5, 1.0, '4')



```
In [23]: #importing tensorflow
from tensorflow.keras.layers import Dense, Input, Conv2D, MaxPool2D, Activation, Dropout, Flatten
from tensorflow.keras.models import Model, Sequential
import random as rn
```

In [23]:

Image Data Generator

Before doing Image Data Generator let's split the dataset using stratified sampling

```
In [24]: train_df, validation_df = train_test_split(imagelabel,
                                                    test_size=0.2,
                                                    random_state=42,
                                                    shuffle=True,
                                                    stratify = imagelabel['label']
                                                    )
```

```
In [25]: train_df.head()
```

Out[25]:

	path	label
10155	imagesj/j/i/w/jiw43c00/89638854_8861.tif	6
45402	imagesb/b/x/k/bxk05e00/2040785858.tif	10
34244	imagesi/i/u/d/iud42e00/2500090606_2500090631.tif	12
46324	imagesp/p/v/z/pvz46c00/2505161384_1385.tif	1
25642	imagesf/f/q/c/fqc66d00/504315522+-5526.tif	3

```
In [26]: train_df.label.value_counts()
```

```
Out[26]: 0      2413
          14      2405
          13      2405
          12      2405
           3      2404
           8      2402
           9      2402
          10      2402
           7      2400
           5      2399
          15      2397
           1      2395
           4      2395
           2      2394
          11      2394
           6      2388
          Name: label, dtype: int64
```

```
In [27]: validation_df.head()
```

```
Out[27]:
```

	path	label
8538	imagesx/x/p/d/xpd30a00/60004171_60004175.tif	6
2381	imagesd/d/z/y/dzy35e00/2040765737_2040765738.tif	10
15246	imagese/e/k/q/ekq43f00/0030049206.tif	4
32549	imagesc/c/m/q/cmqq21a00/0071019330.tif	13
41562	imagesq/q/d/i/qdi71a00/2057432757_2057432767.tif	7


```
In [28]: validation_df.label.value_counts()
```

```
Out[28]: 0      603  
        13     602  
         8     601  
         3     601  
        14     601  
        12     601  
         9     600  
         5     600  
         7     600  
        10     600  
        15     599  
         1     599  
         4     599  
         2     599  
        11     598  
         6     597  
        Name: label, dtype: int64
```

```
In [28]:
```

```
In [29]: imagelabel.iloc[33416] # we are cross-checking our dataset
```

```
Out[29]: path      imagesd/d/u/f/duf23c00/96313362.tif  
        label      4  
        Name: 33416, dtype: object
```

```
In [30]: h,w = 224, 224 # image should be resized
```

```

In [31]: ### Image data generator class
datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rescale=1./255,
    rotation_range=10,
    shear_range=0.3,
    zoom_range=0.2
)
test_datagen = tf.keras.preprocessing.image.ImageDataGenerator(
    rescale=1./255
)

train_dataset = datagen.flow_from_dataframe(
    dataframe=train_df,
    x_col='path',
    y_col='label',
    directory=dir_path,
    target_size=(h,w),
    seed=42,
    batch_size=32, # use small batch size to overcome memory warning
    # subset="training",
    class_mode="sparse"
)
validation_dataset = test_datagen.flow_from_dataframe(
    dataframe=validation_df,
    x_col='path',
    y_col='label',
    directory=dir_path,
    target_size=(h,w),
    seed=42,
    batch_size=32,
    # subset="validation",
    class_mode="sparse"
)

```

Found 38400 validated image filenames belonging to 16 classes.
Found 9600 validated image filenames belonging to 16 classes.

```

In [32]: n_train_sample = int(imagelabel.shape[0]*(1-0.2))
n_val_sample = int(imagelabel.shape[0]*0.2)
batch_size=32
n_train_sample, n_val_sample

```

Out[32]: (38400, 9600)

Also include callbacks

```
In [33]: # save model if accuracy is improved
from tensorflow.keras.callbacks import ModelCheckpoint
filepathw = 'model_save/weights-{epoch:02d}-{val_accuracy:.4f}.hdf5'

checkpoint = ModelCheckpoint(
    filepath = filepathw,
    monitor = 'val_accuracy',
    verbose = 1,
    save_best_only=True,
    mode = 'max' # max in case of accuracy # min=> loss # auto=> detect automatic
```

```
In [34]: # decay Learning rate
class decayLR(tf.keras.callbacks.Callback):
    """
    This is custom callback class to implement decay the Learning Rate.
    Author: Muhammad Iqbal Bajmi @AppliedAICourse.com
    """
    def on_train_begin(self, logs={}):
        self.valid_acc = {'accuracy': []}

    def on_epoch_begin(self, epoch, logs={}):
        print(epoch)
        print("Learning rate is : {}".format(float(self.model.optimizer.learning_rate)))
        if epoch>1:
            if (epoch+1)%3==0:
                lr = self.model.optimizer.learning_rate
                self.model.optimizer.learning_rate = self.model.optimizer.learning_rate - (lr*0.05)
    def on_epoch_end(self, epoch, logs={}):
        self.valid_acc['accuracy'].append(logs.get('val_accuracy'))
        if epoch>0: # because epoch starts from 0
            if self.valid_acc['accuracy'][epoch] < self.valid_acc['accuracy'][epoch-1]:
                lr = self.model.optimizer.learning_rate
                self.model.optimizer.learning_rate = self.model.optimizer.learning_rate - (lr*0.1)

decaylr = decayLR()
```

```
In [35]: # TerminateOnNaN from official tensorflow
from tensorflow.keras.callbacks import TerminateOnNaN
terminate = TerminateOnNaN()
```

```
In [36]: # Terminateon Nan
class TerminateNaN(tf.keras.callbacks.Callback):
    def __init__(self):
        self.we = 0
    def on_epoch_end(self, epoch, logs={}):
        import numpy as np
        loss = logs.get('loss')
        w = self.model.get_weights()
        w = np.array(w)
        self.we = w
        if loss is not None:
            if np.isnan(loss) or np.isinf(loss):
                print("Invalid loss and terminated at epoch {}".format(epoch))
                self.model.stop_training = True
        no_of_layers = len(self.model.layers)
        for i in range(no_of_layers):
            layer = self.model.layers[i]
            weights = layer.get_weights()[0]
            biases = layer.get_weights()[1]

            if np.isnan(weights).any() or np.isinf(weights).any():
                print("Invalid weights and terminated at epoch {}".format(epoch))
                self.model.stop_training = True

terminate_nan = TerminateNaN()
```

```
In [37]: # Stop training if validation_accuracy not improved
from tensorflow.keras.callbacks import EarlyStopping

earlystop = EarlyStopping(
    monitor = 'val_accuracy',
    mode = 'max',
    patience = 2 # wait for 2 epochs, terminate if no improvement
)
```

```
In [38]: import datetime
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=1, write_graph=1)
```

```
In [39]: callback_list = [checkpoint, decaylr, terminate, earlystop, tensorboard_callback]
```

Model-1

1. Use [VGG-16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) (https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling), 2 FC layers and a output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer**
4. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

Adding Layers to the model and then train

In [38]:

```
In [47]: # clear all previous sessions
tf.keras.backend.clear_session()
# initializing sequential model
model_1 = tf.keras.models.Sequential()
# Load VGG16
model = tf.keras.applications.VGG16(
    include_top = False,
    weights = 'imagenet',
    input_shape=(h,w,3)
)
# making all layers.trainable False
for layer in model.layers:
    layer.trainable = False

model_1.add(model) # adding VGG model to our Sequential model
model_1.add(Conv2D(filters=256,
    kernel_size=(3,3),
    strides = (1,1),
    padding='valid',
    data_format = 'channels_last',
    activation='relu',
    kernel_initializer=tf.keras.initializers.he_normal(seed=0),
    name = 'Conv1'
    #input_shape = model.output_shape[1:]
))
model_1.add(MaxPool2D(pool_size=(2,2),
    strides=(2,2),
    padding='valid',
    data_format='channels_last',
    name='Pool1'
))
model_1.add(Flatten(data_format='channels_last', name='Flatten'))
model_1.add(Dense(units=256,
    activation='relu',
    kernel_initializer=tf.keras.initializers.glorot_normal(seed=32), name='FC1'
))
model_1.add(Dense(units=128,
    activation='relu',
    kernel_initializer=tf.keras.initializers.glorot_normal(seed=32), name='FC2'
))
model_1.add(Dense(units=64,
    activation='relu',
    kernel_initializer=tf.keras.initializers.glorot_normal(seed=33), name='FC3'
))
model_1.add(Dense(units=16,
    activation='softmax',
    kernel_initializer=tf.keras.initializers.glorot_normal(seed=3), name='Output')
```

```
))
```

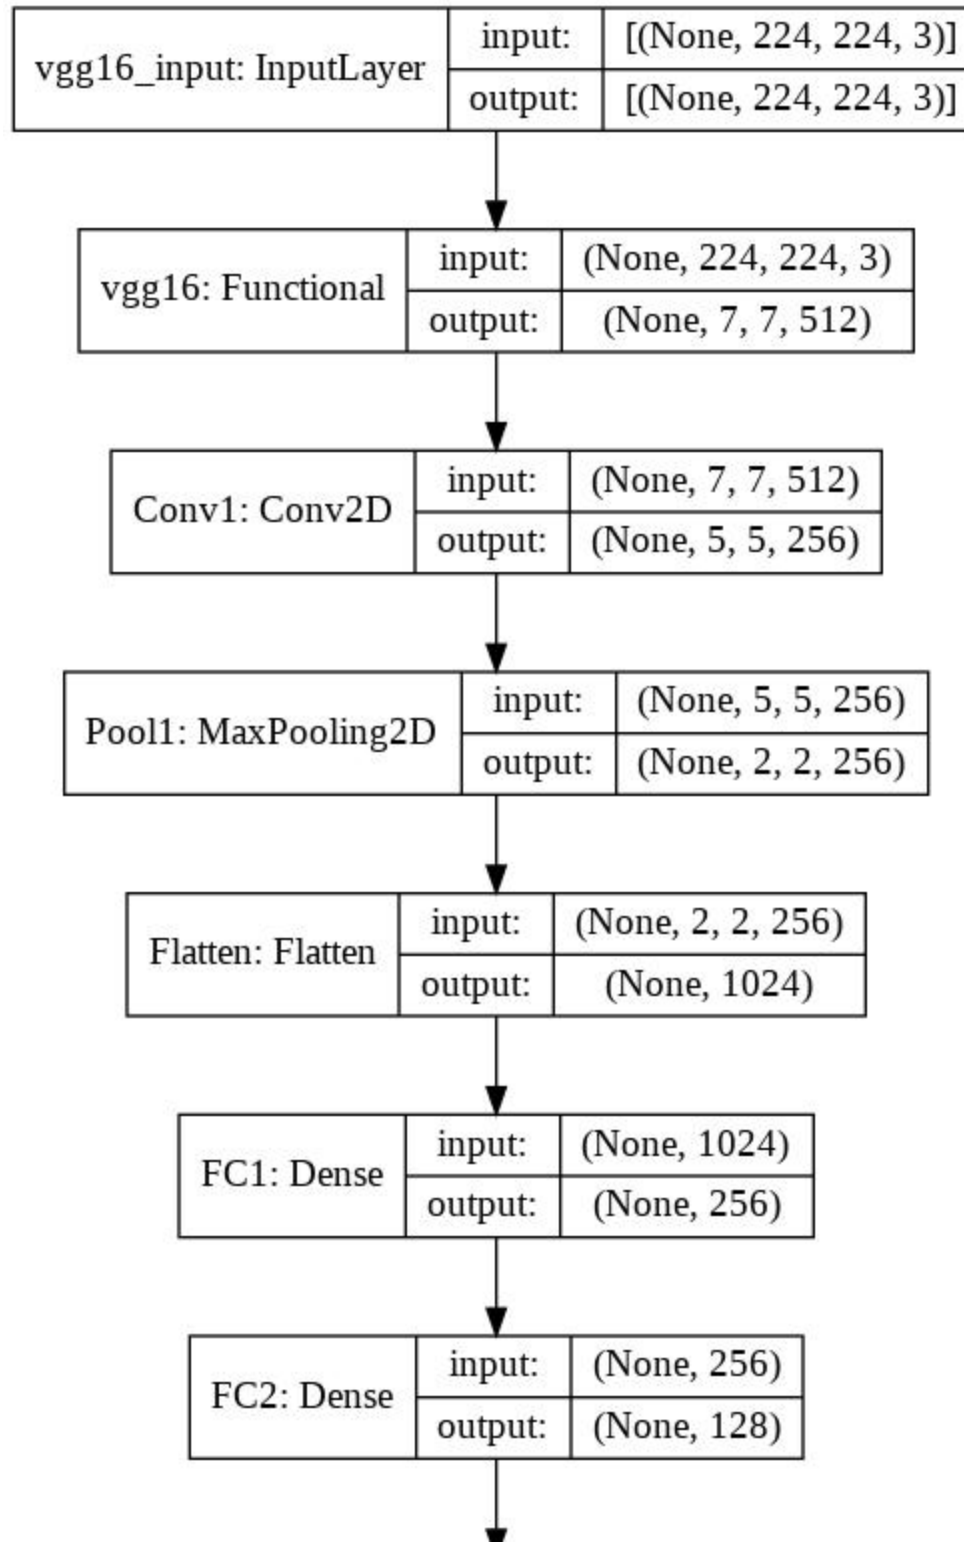
```
In [48]: model_1.summary()
```

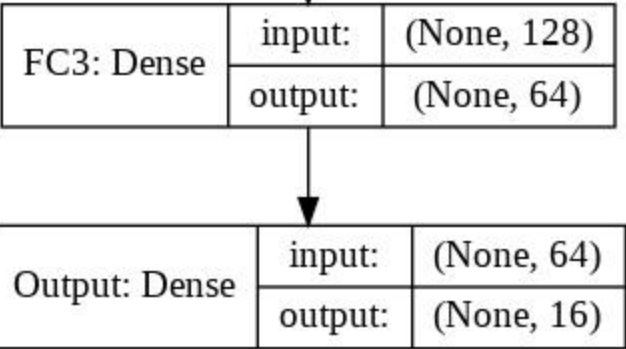
```
Model: "sequential"
```

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 7, 7, 512)	14714688
Conv1 (Conv2D)	(None, 5, 5, 256)	1179904
Pool1 (MaxPooling2D)	(None, 2, 2, 256)	0
Flatten (Flatten)	(None, 1024)	0
FC1 (Dense)	(None, 256)	262400
FC2 (Dense)	(None, 128)	32896
FC3 (Dense)	(None, 64)	8256
Output (Dense)	(None, 16)	1040
=====		
Total params: 16,199,184		
Trainable params: 1,484,496		
Non-trainable params: 14,714,688		
=====		

```
In [49]: tf.keras.utils.plot_model(model_1, to_file='model1.jpeg', show_shapes=True)
```

```
Out[49]:
```





```
In [51]: # remove all previous logs
!rm -rf logs
# remove all saved models
!rm -rf model_save
#compile the model
model_1.compile(optimizer=tf.keras.optimizers.Adam(),
                loss = 'sparse_categorical_crossentropy',
                # loss=tf.keras.losses.sparse_categorical_crossentropy(), # use sparse categorical cross entropy for non
                # loss='categorical_crossentropy', # use categorical cross entropy for one-hot encoded labels
                metrics=['accuracy'])

# fit the model
model_1.fit(x = train_dataset,
           validation_data=validation_dataset,
           epochs=10,
           callbacks=[callback_list]
           )
```

Epoch 1/10

0

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 595s 495ms/step - loss: 1.6186 - accuracy: 0.4966 - val_loss: 1.2469 - val_accuracy: 0.6051

Epoch 00001: val_accuracy improved from -inf to 0.60510, saving model to model_save/weights-01-0.6051.hdf5

Epoch 2/10

1

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 597s 497ms/step - loss: 1.2189 - accuracy: 0.6247 - val_loss: 1.1454 - val_accuracy: 0.6468

Epoch 00002: val_accuracy improved from 0.60510 to 0.64677, saving model to model_save/weights-02-0.6468.hdf5

Epoch 3/10

2

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 583s 486ms/step - loss: 1.0793 - accuracy: 0.6675 - val_loss: 1.0993 - val_accuracy: 0.6648

Epoch 00003: val_accuracy improved from 0.64677 to 0.66479, saving model to model_save/weights-03-0.6648.hdf5

Epoch 4/10

3

Learning rate is : 0.0009500000160187483

1200/1200 [=====] - 575s 479ms/step - loss: 0.9874 - accuracy: 0.6945 - val_loss: 1.0194 - val_accuracy: 0.6974

Epoch 00004: val_accuracy improved from 0.66479 to 0.69740, saving model to model_save/weights-04-0.6974.hdf5

```
Epoch 5/10
4
Learning rate is : 0.0009500000160187483
1200/1200 [=====] - 576s 480ms/step - loss: 0.9304 - accuracy: 0.7140 - val_loss: 1.0042 - va
l_accuracy: 0.7006

Epoch 00005: val_accuracy improved from 0.69740 to 0.70063, saving model to model_save/weights-05-0.7006.hdf5
Epoch 6/10
5
Learning rate is : 0.0009500000160187483
1200/1200 [=====] - 569s 474ms/step - loss: 0.8877 - accuracy: 0.7273 - val_loss: 0.9920 - va
l_accuracy: 0.7134

Epoch 00006: val_accuracy improved from 0.70063 to 0.71344, saving model to model_save/weights-06-0.7134.hdf5
Epoch 7/10
6
Learning rate is : 0.0009025000035762787
1200/1200 [=====] - 558s 465ms/step - loss: 0.8421 - accuracy: 0.7407 - val_loss: 0.9713 - va
l_accuracy: 0.7159

Epoch 00007: val_accuracy improved from 0.71344 to 0.71594, saving model to model_save/weights-07-0.7159.hdf5
Epoch 8/10
7
Learning rate is : 0.0009025000035762787
1200/1200 [=====] - 546s 455ms/step - loss: 0.8108 - accuracy: 0.7478 - val_loss: 1.0405 - va
l_accuracy: 0.6938

Epoch 00008: val_accuracy did not improve from 0.71594
Epoch 9/10
8
Learning rate is : 0.000812250014860183
1200/1200 [=====] - 548s 457ms/step - loss: 0.7658 - accuracy: 0.7648 - val_loss: 0.9602 - va
l_accuracy: 0.7210

Epoch 00009: val_accuracy improved from 0.71594 to 0.72104, saving model to model_save/weights-09-0.7210.hdf5
Epoch 10/10
9
Learning rate is : 0.0007716375403106213
1200/1200 [=====] - 546s 455ms/step - loss: 0.7304 - accuracy: 0.7759 - val_loss: 0.9391 - va
l_accuracy: 0.7328

Epoch 00010: val_accuracy improved from 0.72104 to 0.73281, saving model to model_save/weights-10-0.7328.hdf5
```

Out[51]: <tensorflow.python.keras.callbacks.History at 0x7f827c118dd0>

Tensorboard screenshot of Model-1



+ Code + Text

✓ RAM
Disk

Editing



☒ Ignore outliers in chart scaling

Tooltip sorting method: default

Smoothing

0.6

Horizontal Axis

STEP RELATIVE WALL

Runs

Write a regex to filter runs


☒ 20210322-103502/train

☒ 20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

epoch_accuracy



Automatic saving failed. This file was updated remotely or in another tab.

[Show diff](#)

	Smoothed	Value	Step	Time	Relative
rain	0.7548	0.773	9	Mon Mar 22, 18:07:21	1h 25m 9s
20210322-103502/validation	0.7183	0.7328	9	Mon Mar 22, 18:07:21	1h 25m 9s

epoch_loss





+ Code + Text

✓ RAM
Disk

Editing



Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs



20210322-103502/train




20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

Name	Smoothed	Value	Step	Time	Relative
20210322-103502/train	0.8801	0.8201	7	Mon Mar 22, 17:49:04	1h 6m 52s
20210322-103502/validation	1.018	1.04	7	Mon Mar 22, 17:49:04	1h 6m 52s

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

[]



[]



+ Code + Text

✓ RAM
Disk

Editing



Runs

Write a regex to filter runs



20210322-103502/train



20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

STEP

RELATIVE

WALL

Conv1



2 ^

Conv1/bias_0

20210322-103502/train

Conv1/kernel_0

20210322-103502/train



FC1



2 ^

FC1/bias_0

20210322-103502/train

FC1/kernel_0

20210322-103502/train





+ Code + Text

✓ RAM
Disk

Editing



Write a regex to filter runs

☒ 20210322-103502/train☒ 20210322-103502/validation

TOGGLE ALL RUNS


logs/fit

FC2

2 ^


FC2/bias_0

20210322-103502/train



FC2/kernel_0

20210322-103502/train




FC3

2 ^


FC3/bias_0


20210322-103502/train



FC3/kernel_0

20210322-103502/train





+ Code + Text

✓ RAM
Disk

Editing



Write a regex to filter runs

☒ 20210322-103502/train☒ 20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

Output


2 ^

Output/bias_0

20210322-103502/train

Output/kernel_0

20210322-103502/train



block1_conv1



2 ^

block1_conv1/bias_0

20210322-103502/train

block1_conv1/kernel_0

20210322-103502/train





+ Code + Text

✓ RAM
Disk

Editing



OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210322-103502/train☒ 20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

Conv1



2 ^

Conv1/bias_0

20210322-103502/train

Conv1/kernel_0

20210322-103502/train



FC1



2 ^

FC1/bias_0

20210322-103502/train

FC1/kernel_0

20210322-103502/train





+ Code + Text

✓ RAM
Disk

Editing



OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210322-103502/train☒ 20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

FC2



2 ^

FC2/bias_0

20210322-103502/train

FC2/kernel_0

20210322-103502/train



FC3



2 ^

FC3/bias_0

20210322-103502/train

FC3/kernel_0

20210322-103502/train





+ Code + Text



OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210322-103502/train☒ 20210322-103502/validation

TOGGLE ALL RUNS

logs/fit

✓ RAM
Disk

Editing



Output



2 ^

Output/bias_0

20210322-103502/train

Output/kernel_0

20210322-103502/train



block1_conv1



2 ^

block1_conv1/bias_0

20210322-103502/train

block1_conv1/kernel_0

20210322-103502/train



In [51]:

In [52]: %load_ext tensorboard

In [53]: %tensorboard --logdir logs/fit

<IPython.core.display.Javascript object>

In [44]:

In [44]:

In [44]:

Model-2

1. Use [VGG-16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) (https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size 7×7×512 can be equivalently expressed as a CONV layer with F=7,P=0,S=1,K=4096. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be 1×1×4096 since only a single depth column “fits” across the input volume, giving identical result as the initial FC layer. You can refer [this](http://cs231n.github.io/convolutional-networks/#convert) (<http://cs231n.github.io/convolutional-networks/#convert>) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**
3. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

In [84]:

```
# clear all previous sessions
tf.keras.backend.clear_session()
# initializing sequential model
model_2 = tf.keras.models.Sequential()
# Load VGG16
model = tf.keras.applications.VGG16(
    include_top = False,
    weights = 'imagenet',
    input_shape=(h,w,3)
)
# making all layers.trainable True
for layer in model.layers:
    layer.trainable = False

model_2.add(model) # adding VGG model to our Sequential model
model_2.add(Conv2D(filters=4096,
    kernel_size=(7,7),
    strides = (1,1),
    padding='valid', # valid means no padding
    data_format = 'channels_last',
    activation='relu',
    kernel_initializer=tf.keras.initializers.he_uniform(seed=42),
    name = 'fc_Conv1'
    #input_shape = model.output_shape[1:]
))
model_2.add(Conv2D(
    filters=4096,
    kernel_size=(1,1),
    padding='valid',
    data_format='channels_last',
    activation='relu',
    kernel_initializer=tf.keras.initializers.he_uniform(seed=42),
    name='fc_conv2'
))
model_2.add(Flatten(data_format='channels_last',name='Flatten_last'))
model_2.add(Dense(units=16,
    activation='softmax',
    kernel_initializer=tf.keras.initializers.glorot_normal(seed=3), name='Output'
))
```

In [85]:

model_2.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 7, 7, 512)	14714688

fc_Conv1 (Conv2D)	(None, 1, 1, 4096)	102764544

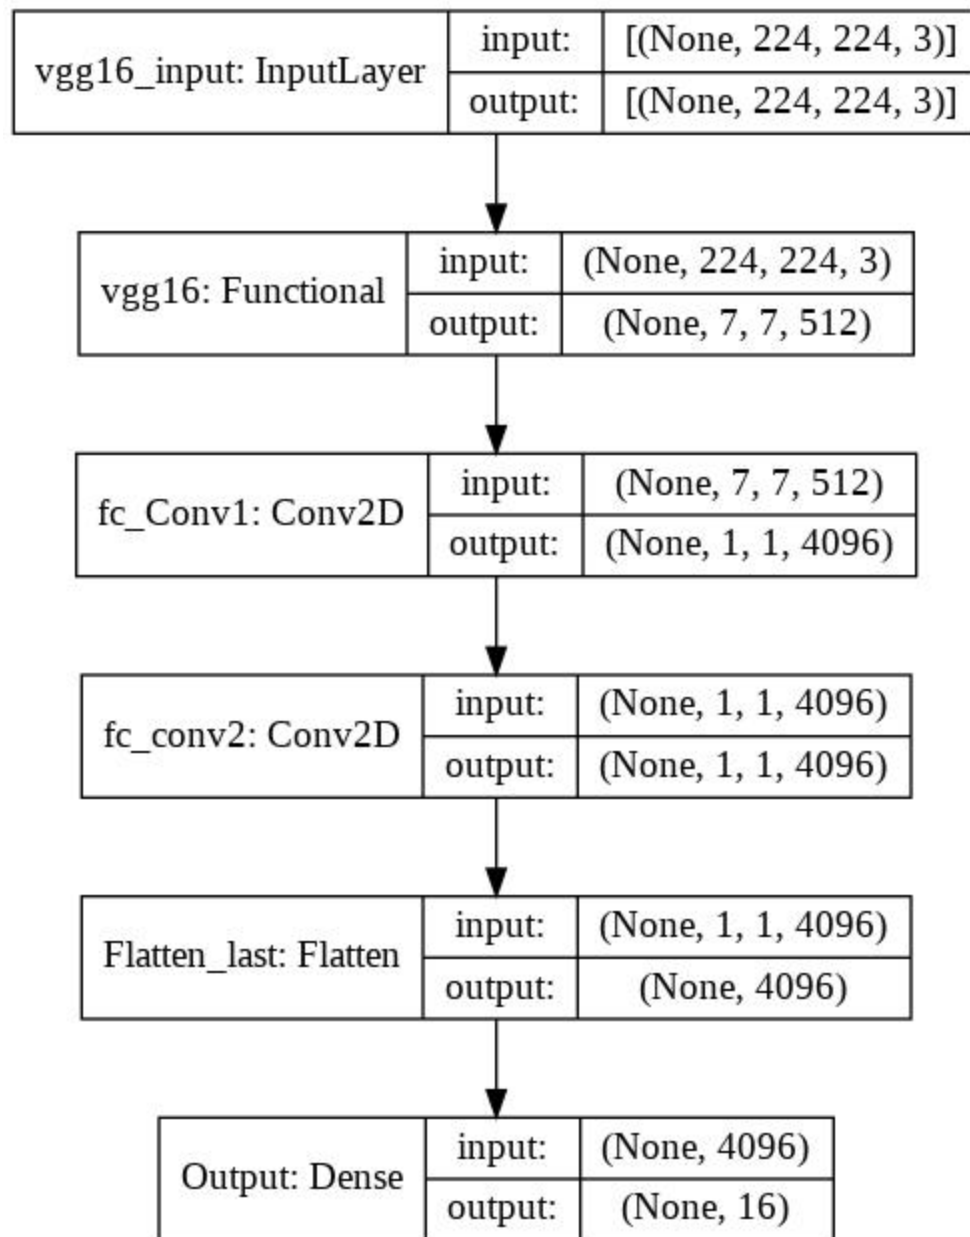
fc_conv2 (Conv2D)	(None, 1, 1, 4096)	16781312

Flatten_last (Flatten)	(None, 4096)	0

Output (Dense)	(None, 16)	65552
=====		
Total params: 134,326,096		
Trainable params: 119,611,408		
Non-trainable params: 14,714,688		

In [86]: `tf.keras.utils.plot_model(model_2, to_file='model1.jpeg', show_shapes=True)`

Out[86]:



```
In [87]: # train_df.head(32)
```



```
In [88]: # remove previous logs
!rm -rf logs
# remove previous saved models
!rm -rf model_save
#compile the model
model_2.compile(optimizer=tf.keras.optimizers.Adam(),
                loss='sparse_categorical_crossentropy',
                metrics=['accuracy'])
#fit the model
model_2.fit(x = train_dataset,
            validation_data=validation_dataset,
            epochs=10,
            callbacks=[callback_list]
            )
```

Epoch 1/10

0

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 666s 554ms/step - loss: 2.6802 - accuracy: 0.4623 - val_loss: 1.1420 - val_accuracy: 0.6556

Epoch 00001: val_accuracy improved from -inf to 0.65562, saving model to model_save/weights-01-0.6556.hdf5

Epoch 2/10

1

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 677s 564ms/step - loss: 1.1569 - accuracy: 0.6468 - val_loss: 1.0220 - val_accuracy: 0.6881

Epoch 00002: val_accuracy improved from 0.65562 to 0.68813, saving model to model_save/weights-02-0.6881.hdf5

Epoch 3/10

2

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 671s 559ms/step - loss: 1.0431 - accuracy: 0.6815 - val_loss: 0.9893 - val_accuracy: 0.7019

Epoch 00003: val_accuracy improved from 0.68813 to 0.70187, saving model to model_save/weights-03-0.7019.hdf5

Epoch 4/10

3

Learning rate is : 0.0009500000160187483

1200/1200 [=====] - 674s 561ms/step - loss: 0.9872 - accuracy: 0.6946 - val_loss: 0.9673 - val_accuracy: 0.7061

Epoch 00004: val_accuracy improved from 0.70187 to 0.70615, saving model to model_save/weights-04-0.7061.hdf5

Epoch 5/10

4

Learning rate is : 0.0009500000160187483

1200/1200 [=====] - 674s 561ms/step - loss: 0.9195 - accuracy: 0.7198 - val_loss: 0.9786 - val_

accuracy: 0.7063

Epoch 00005: val_accuracy improved from 0.70615 to 0.70625, saving model to model_save/weights-05-0.7063.hdf5

Epoch 6/10

5

Learning rate is : 0.0009500000160187483

1200/1200 [=====] - 675s 562ms/step - loss: 0.8758 - accuracy: 0.7335 - val_loss: 0.9532 - val_accuracy: 0.7200

Epoch 00006: val_accuracy improved from 0.70625 to 0.72000, saving model to model_save/weights-06-0.7200.hdf5

Epoch 7/10

6

Learning rate is : 0.0009025000035762787

1200/1200 [=====] - 665s 554ms/step - loss: 0.8424 - accuracy: 0.7395 - val_loss: 0.9762 - val_accuracy: 0.7138

Epoch 00007: val_accuracy did not improve from 0.72000

Epoch 8/10

7

Learning rate is : 0.000812250014860183

1200/1200 [=====] - 648s 540ms/step - loss: 0.7893 - accuracy: 0.7574 - val_loss: 0.9574 - val_accuracy: 0.7144

Epoch 00008: val_accuracy did not improve from 0.72000

Out[88]: <tensorflow.python.keras.callbacks.History at 0x7ff281687750>

Tensorboard screenshots of Model-2



+ Code + Text

✓ RAM
Disk

Editing



Tooltip sorting method: default

Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs


☒ 20210323-013000/train☒ 20210323-013000/validation

TOGGLE ALL RUNS


logs/fit



epoch_accuracy



epoch_loss

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



+ Code + Text

✓ RAM
Disk

Editing



Horizontal axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210323-013000/train☒ 20210323-013000/validation

TOGGLE ALL RUNS


logs/fit

Output

2 ^


Output/bias_0

20210323-013000/train



Output/kernel_0

20210323-013000/train




block1_conv1

2 ^


block1_conv1/bias_0

20210323-013000/train



block1_conv1/kernel_0

20210323-013000/train

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



+ Code + Text

✓ RAM
Disk

Editing



Write a regex to filter runs


☒ 20210323-013000/train☒ 20210323-013000/validation

TOGGLE ALL RUNS

logs/fit


fc_Conv1/bias_0

20210323-013000/train



fc_Conv1/kernel_0

20210323-013000/train




fc_conv2

2 ^


fc_conv2/bias_0

20210323-013000/train



fc_conv2/kernel_0

20210323-013000/train

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



+ Code + Text

✓ RAM
Disk

Editing



Histogram mode

OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210323-013000/train☒ 20210323-013000/validation

TOGGLE ALL RUNS

logs/fit



Output



2 ^

Output/bias_0

20210323-013000/train

Output/kernel_0

20210323-013000/train



block1_conv1



2 ^

block1_conv1/bias_0

20210323-013000/train

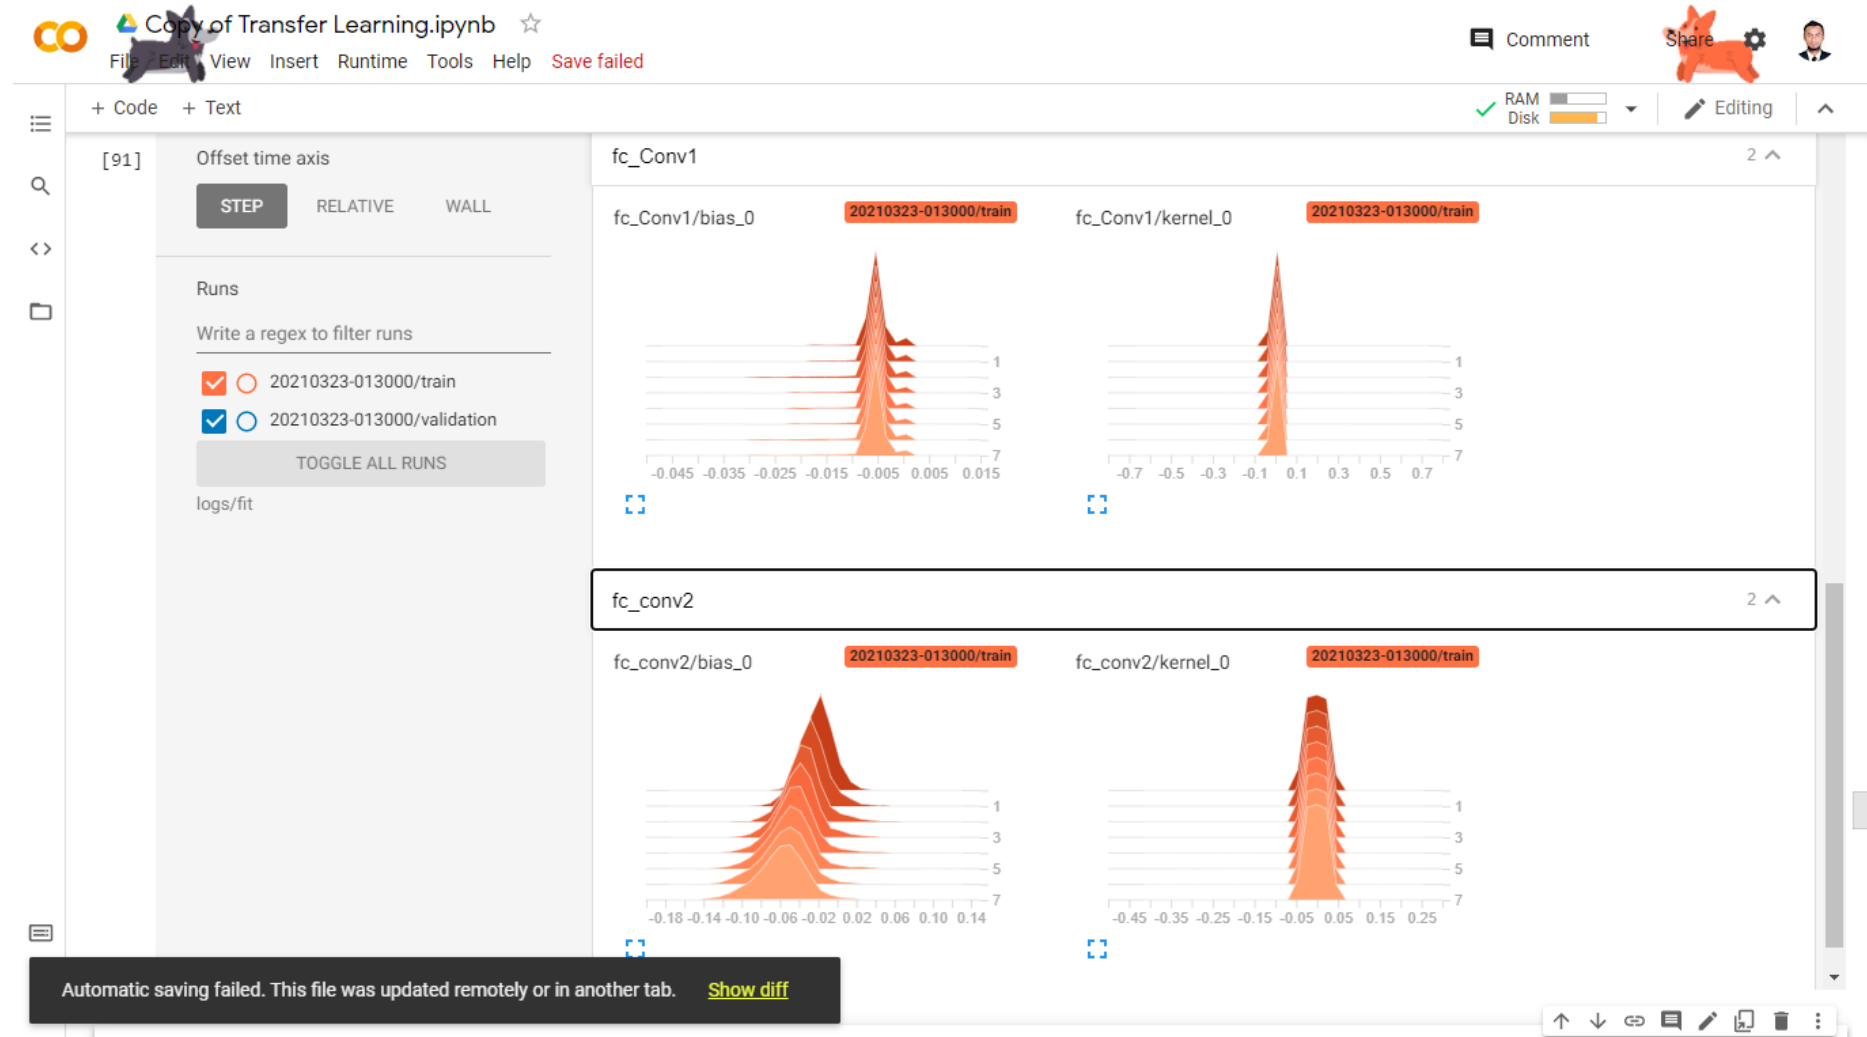
block1_conv1/kernel_0

20210323-013000/train



Automatic saving failed. This file was updated remotely or in another tab.

[Show diff](#)



In []:

In [90]: `%load_ext tensorboard`

In [91]: `%tensorboard --logdir logs/fit`

<IPython.core.display.Javascript object>

Model-3

1. Use same network as Model-2 'INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Out

put Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.


```
In [120]: model = tf.keras.applications.VGG16(
            include_top = False,
            weights = 'imagenet',
            input_shape=(h,w,3)
        )
# making last 6 layers.trainable True
print(len(model.layers)-5)
for i in range(len(model.layers)-5):
    layer = model.layers[i]
    layer.trainable = False
model.summary()
```

14

Model: "vgg16"

Layer (type)	Output Shape	Param #
=====		
input_3 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0

block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
=====		
Total params: 14,714,688		
Trainable params: 7,079,424		
Non-trainable params: 7,635,264		

In [120]:

In [120]:

```
In [92]: # clear all previous sessions
tf.keras.backend.clear_session()
# initializing sequential model
top_model = tf.keras.models.Sequential()
# Load VGG16
model = tf.keras.applications.VGG16(
    include_top = False,
    weights = 'imagenet',
    input_shape=(h,w,3)
)
# making last 6 layers.trainable True
for i in range(len(model.layers)-5):
    layer = model.layers[i]
    layer.trainable = False

top_model.add(model) # adding VGG model to our Sequential model
top_model.add(Conv2D(filters=4096,
    kernel_size=(7,7),
    strides = (1,1),
    padding='valid', # valid means no padding
    data_format = 'channels_last',
    activation=tf.keras.layers.LeakyReLU(),
    kernel_initializer=tf.keras.initializers.he_normal(seed=0),
    name = 'fc_Conv1'
    #input_shape = model.output_shape[1:]
))
top_model.add(Conv2D(
    filters=4096,
    kernel_size=(1,1),
    padding='valid',
    data_format='channels_last',
    activation=tf.keras.layers.LeakyReLU(),
    kernel_initializer=tf.keras.initializers.he_normal(seed=0),
    name='fc_conv2'
))
top_model.add(Flatten(data_format='channels_last',name='Flatten_last'))
top_model.add(Dense(units=16,
    activation='softmax',
    kernel_initializer=tf.keras.initializers.glorot_normal(seed=3), name='Output'
))
```

In [93]: top_model.summary()

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
vgg16 (Functional)	(None, 7, 7, 512)	14714688

fc_Conv1 (Conv2D)	(None, 1, 1, 4096)	102764544

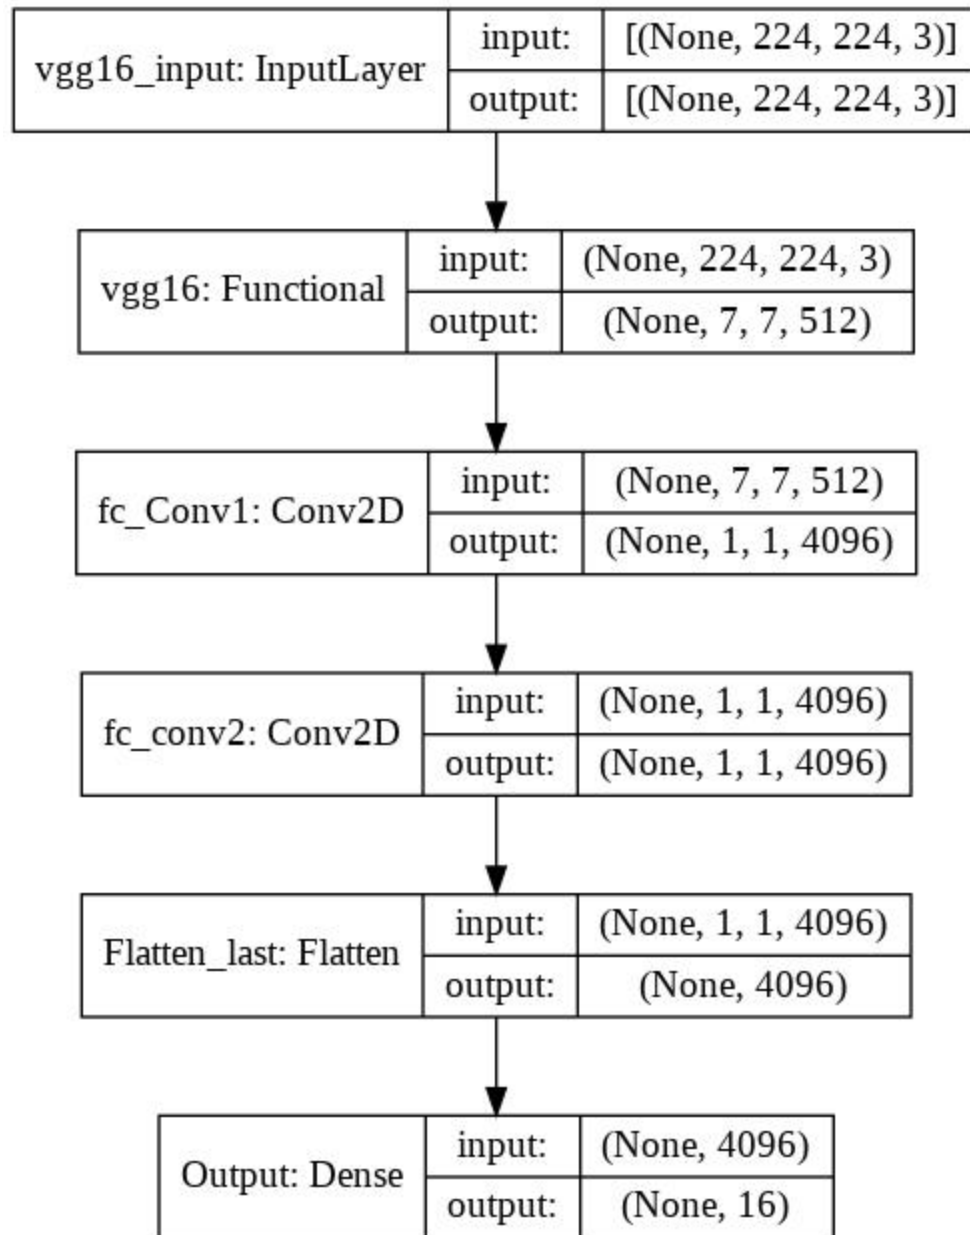
fc_conv2 (Conv2D)	(None, 1, 1, 4096)	16781312

Flatten_last (Flatten)	(None, 4096)	0

Output (Dense)	(None, 16)	65552
=====		
Total params: 134,326,096		
Trainable params: 126,690,832		
Non-trainable params: 7,635,264		

In [95]: `tf.keras.utils.plot_model(top_model, to_file='model3.jpeg', show_shapes=True)`

Out[95]:




```
In [97]: # remove previous logs
!rm -rf logs
# remove previous saved models
!rm -rf model_save
#compile the model
top_model.compile(optimizer=tf.keras.optimizers.Adam(),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

#fit the model
top_model.fit(x = train_dataset,
              validation_data=validation_dataset,
              epochs=10,
              callbacks=[callback_list]
              )
```

Epoch 1/10

0

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 736s 612ms/step - loss: 1295.5925 - accuracy: 0.0628 - val_loss: 165.3107 - val_accuracy: 0.0624

Epoch 00001: val_accuracy did not improve from 0.72000

Epoch 2/10

1

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 723s 602ms/step - loss: 192.8512 - accuracy: 0.0612 - val_loss: 168.8053 - val_accuracy: 0.0625

Epoch 00002: val_accuracy did not improve from 0.72000

Epoch 3/10

2

Learning rate is : 0.0010000000474974513

1200/1200 [=====] - 721s 600ms/step - loss: 174.8353 - accuracy: 0.0615 - val_loss: 172.5757 - val_accuracy: 0.0623

Epoch 00003: val_accuracy did not improve from 0.72000

Epoch 4/10

3

Learning rate is : 0.0008549999911338091

1200/1200 [=====] - 717s 597ms/step - loss: 182.1918 - accuracy: 0.0598 - val_loss: 104.1381 - val_accuracy: 0.0624

Epoch 00004: val_accuracy did not improve from 0.72000

Out[97]: <tensorflow.python.keras.callbacks.History at 0x7ff2204ba790>

Tensorboard screenshots of Model-3



Copy of Transfer Learning.ipynb ☆

File Edit View Insert Runtime Tools Help Save failed

Comment

Share



+ Code + Text



Tooltip sorting method: default



Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs


☒ 20210323-013000/train

☒ 20210323-013000/validation

TOGGLE ALL RUNS


logs/fit

epoch_accuracy



epoch_loss

epoch_loss



Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



+ Code + Text



OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210323-013000/train☒ 20210323-013000/validation

TOGGLE ALL RUNS

logs/fit

✓ RAM
Disk

Editing



Output



2 ^

Output/bias_0

20210323-013000/train

Output/kernel_0

20210323-013000/train



block1_conv1



2 ^

block1_conv1/bias_0

20210323-013000/train

block1_conv1/kernel_0

20210323-013000/train

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)



+ Code + Text



OVERLAY

OFFSET

Offset time axis

STEP

RELATIVE

WALL

Runs

Write a regex to filter runs

☒ 20210323-013000/train☒ 20210323-013000/validation

TOGGLE ALL RUNS

logs/fit



fc_Conv1



2 ^

fc_Conv1/bias_0

20210323-013000/train

fc_Conv1/kernel_0

20210323-013000/train



fc_conv2



2 ^

fc_conv2/bias_0

20210323-013000/train

fc_conv2/kernel_0

20210323-013000/train

Automatic saving failed. This file was updated remotely or in another tab. [Show diff](#)

In []:

In [98]: %tensorboard --logdir logs/fit

<IPython.core.display.Javascript object>

In []:

Result:

Model ----- Accuracy(%)

Model-1 ===== 73.28

Model-2 ===== 72.00

Model-3 ===== 06.25

In []: